



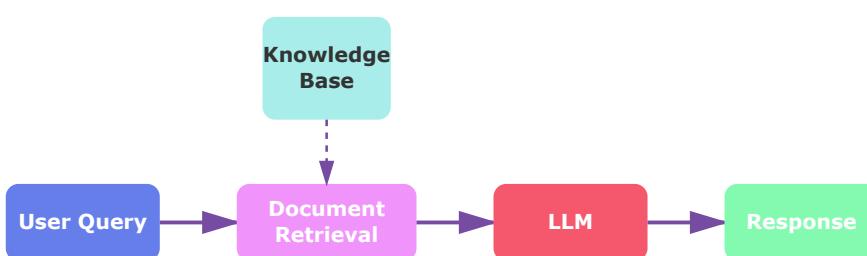
RAG Architectures Guide

Comprehensive Overview of Retrieval-Augmented Generation Systems

What is RAG?

Retrieval-Augmented Generation (RAG) is a technique that enhances Large Language Models (LLMs) by combining their generative capabilities with external knowledge retrieval. Instead of relying solely on the model's training data, RAG systems fetch relevant information from external sources in real-time, leading to more accurate, up-to-date, and contextually relevant responses.

1. Basic/Naive RAG



Overview

The foundational RAG architecture follows a simple three-step process: indexing documents into a vector database, retrieving relevant chunks based on query similarity, and generating responses using the retrieved context.

Key Components

Document Indexing

Documents are chunked and converted into embeddings stored in a vector database.

Similarity Search

Query embeddings are compared with stored embeddings to find relevant documents.

Generation

LLM generates responses using retrieved context and the original query.

Advantages

- Simple to implement and understand
- Fast response times
- Works well for straightforward queries
- Lower computational requirements

Limitations

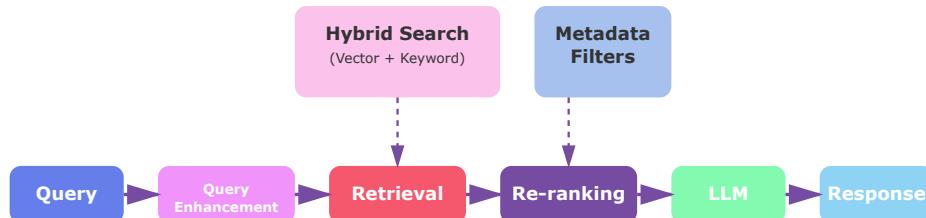
- May retrieve irrelevant chunks
- Limited context window understanding
- Struggles with complex multi-hop queries

- No query refinement mechanism

🎯 Best Use Cases

- Simple FAQ systems
- Document search and summarization
- Basic chatbots with limited knowledge domains

2. 🎯 Advanced RAG



Overview

Advanced RAG improves upon basic RAG by adding pre-retrieval and post-retrieval optimizations. It includes query enhancement, hybrid search strategies, re-ranking mechanisms, and better context management.

Key Enhancements

Query Rewriting

Transforms queries to improve retrieval accuracy using techniques like HyDE or query expansion.

Hybrid Search

Combines vector similarity with keyword-based search for better coverage.

Re-ranking

Uses cross-encoders to re-order retrieved documents by relevance.

Metadata Filtering

Applies filters based on document attributes before or after retrieval.

Advantages

- Higher retrieval accuracy
- Better handling of ambiguous queries
- Improved ranking of relevant documents
- More contextually aware responses

Limitations

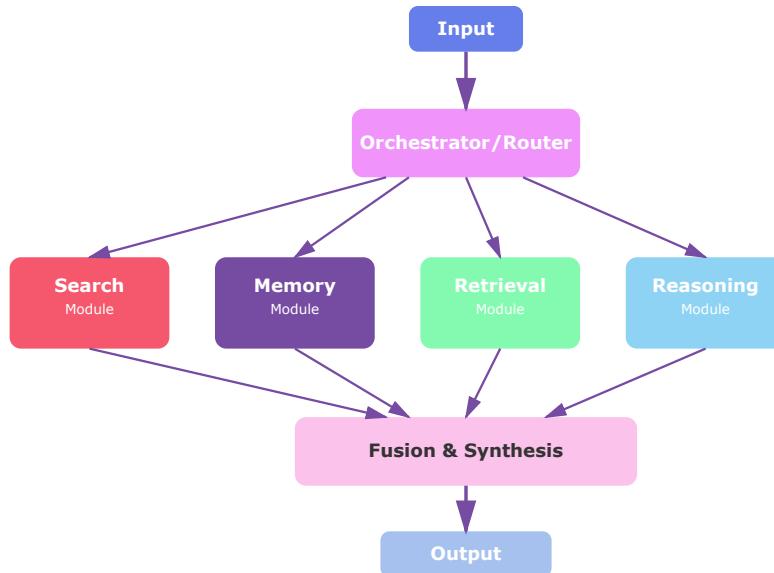
- Increased latency due to multiple steps
- More complex to implement and maintain
- Higher computational costs
- Requires careful tuning of components

Best Use Cases

- Enterprise knowledge management systems

- Legal and medical document search
- Research paper retrieval and analysis
- Customer support with large documentation

3. Modular RAG



Overview

Modular RAG treats the system as a collection of independent, interchangeable modules. It uses a routing mechanism to decide which modules to invoke based on the query type, allowing for flexible and scalable architectures.

Core Modules

 **Router Module**

 **Retrieval Module**

Intelligently directs queries to appropriate processing modules based on intent classification.

Handles document search across multiple knowledge bases with different strategies.

Memory Module

Manages conversation history and contextual information for coherent interactions.

Reasoning Module

Applies logical reasoning and multi-step thinking for complex queries.

Advantages

- Highly flexible and customizable
- Easy to add or remove modules
- Can handle diverse query types
- Better resource optimization
- Easier testing and debugging

Limitations

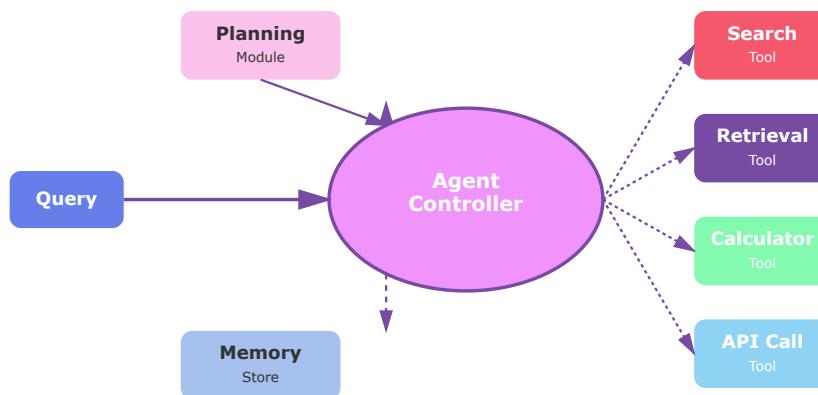
- Complex system design and architecture
- Requires sophisticated routing logic
- Higher development overhead
- Module coordination challenges

Best Use Cases

- Multi-domain knowledge systems
- Enterprise AI assistants with varied capabilities

- Research platforms requiring different reasoning strategies
- Systems needing frequent updates and extensions

4. Agentic RAG



Overview

Agentic RAG uses autonomous agents that can plan, reason, and iteratively use tools to solve complex queries. The agent breaks down problems, decides which tools to use, and synthesizes information from multiple sources.

Key Capabilities

Multi-Step Planning

Agents create execution plans with multiple steps to

Tool Use

Dynamic selection and execution of various tools

tackle complex queries systematically.

including search, calculators, and APIs.

Iterative Refinement

Agents can reflect on results and adjust their approach based on intermediate outcomes.

Memory Management

Maintains context across interactions and learns from past experiences.

Advantages

- Handles complex multi-step reasoning
- Autonomous problem-solving capabilities
- Can use external tools and APIs
- Adaptive to novel situations
- Self-correcting and iterative

Limitations

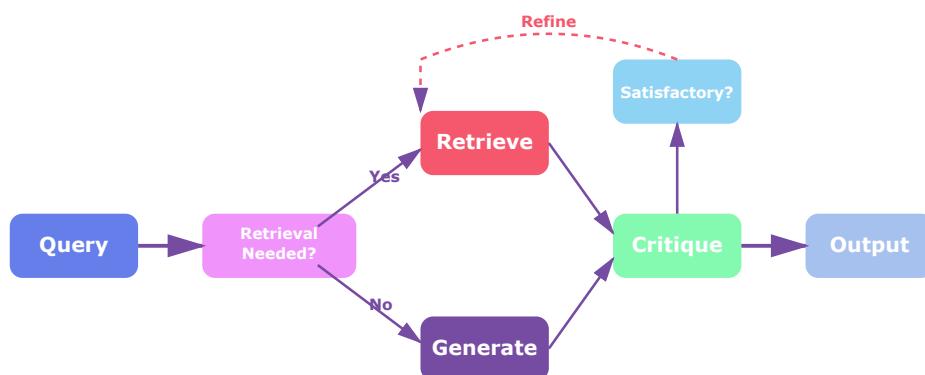
- High computational costs
- Unpredictable behavior in some cases
- Longer response times
- Requires careful safety guardrails
- Complex debugging and monitoring

Best Use Cases

- Research assistants requiring multi-source synthesis
- Complex data analysis and reporting

- Autonomous customer service agents
- Scientific literature review and analysis
- Financial analysis with multiple data sources

5. Self-RAG (Self-Reflective RAG)



Overview

Self-RAG introduces reflection tokens that allow the model to critique its own outputs and decide when retrieval is necessary. It can self-assess relevance, support, and usefulness of retrieved information.

Reflection Mechanisms

 **Retrieval Tokens**

 **Relevance Tokens**

Model decides if external information is needed before retrieving.

Assesses if retrieved documents are relevant to the query.

Support Tokens

Evaluates if the generated response is supported by retrieved evidence.

Utility Tokens

Judges overall usefulness of the final response.

Advantages

- Self-correcting capabilities
- Efficient retrieval (only when needed)
- Better factual accuracy
- Reduced hallucinations
- Transparent reasoning process

Limitations

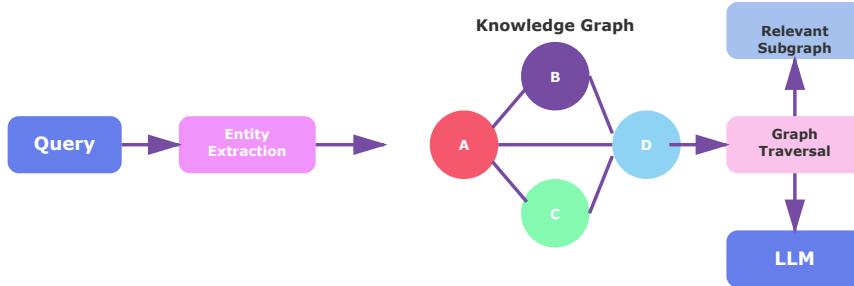
- Requires specially trained models
- Additional tokens increase latency
- More complex evaluation pipeline
- May be overly conservative in retrieval

Best Use Cases

- Fact-checking systems requiring high accuracy
- Medical and legal applications
- Academic research tools

- Systems where transparency is critical

6. GraphRAG



Overview

GraphRAG structures knowledge as a graph where entities are nodes and relationships are edges. It performs graph traversal to find related information, enabling better understanding of relationships and multi-hop reasoning.

Graph Components

Entity Nodes

Represent people, places, concepts, and objects with rich attributes.

Relationship Edges

Typed connections between entities capturing semantic relationships.

Graph Traversal

Navigate through relationships to discover connected information.

Subgraph Extraction

Identify relevant portions of the graph for specific queries.

Advantages

- Excellent for relationship-heavy queries
- Natural multi-hop reasoning
- Rich contextual understanding
- Handles complex entity relationships
- Supports graph algorithms for insights

Limitations

- Complex graph construction and maintenance
- Requires entity extraction and linking
- Computationally expensive for large graphs
- Challenging to keep graphs up-to-date
- May struggle with unstructured information

Best Use Cases

- Knowledge base question answering
- Recommendation systems
- Fraud detection and investigation
- Social network analysis

- Scientific research with interconnected concepts



Architecture Comparison

Architecture	Complexity	Latency	Accuracy	Best For
Basic RAG	Low	Fast	Good	Simple Q&A, FAQs
Advanced RAG	Medium	Medium	Very Good	Enterprise knowledge management
Modular RAG	High	Variable	Excellent	Multi-domain systems
Agentic RAG	Very High	Slow	Excellent	Complex reasoning tasks
Self-RAG	High	Medium-Slow	Excellent	High- accuracy applications
GraphRAG	Very High	Medium	Excellent	Relationship- heavy queries



Key Takeaways

- **Start Simple:** Begin with Basic RAG for straightforward use cases and scale up as needed.
- **Match Architecture to Use Case:** Choose based on query complexity, accuracy requirements, and latency constraints.
- **Consider Trade-offs:** Higher accuracy often comes with increased complexity and latency.
- **Hybrid Approaches:** Many production systems combine multiple architectures for optimal results.
- **Continuous Evaluation:** Regular testing and monitoring are essential regardless of architecture choice.